

# Logistic Regression Report

---

## 1. Decision Boundary using Logistic Regression

We implemented logistic regression from scratch using Python and applied it to a provided dataset. The independent variables (features) were normalized using standardization (mean = 0, std = 1). The model was trained using gradient descent with a learning rate of **0.1** for **1000 iterations**.

**Final Cost after Convergence:** ~0.2569

**Final Weights:** [2.7313, 1.6394]

**Final Bias:** -0.3089

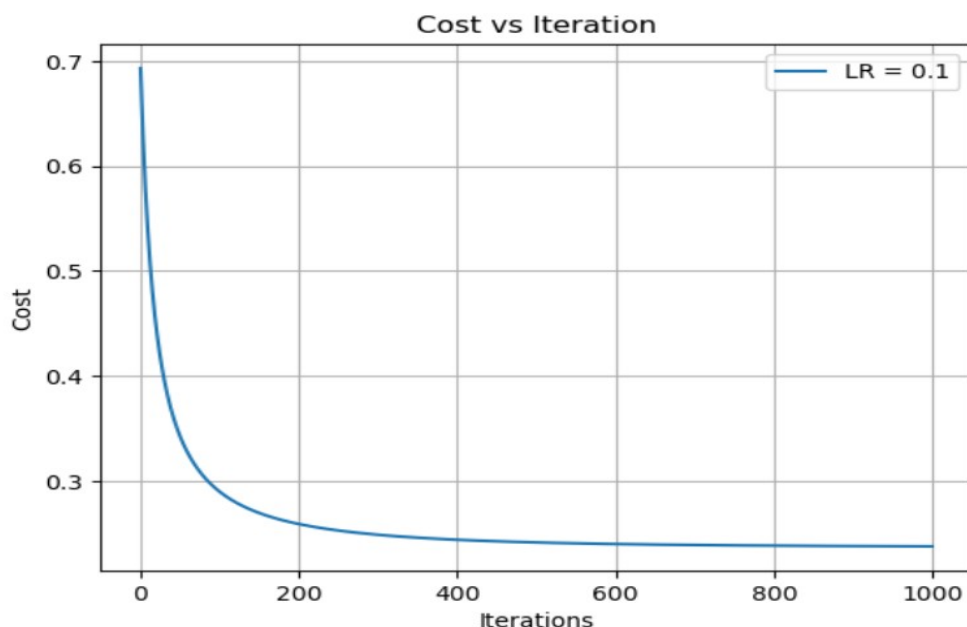
This implies the decision boundary is represented by the equation:

$$W_1X_1 + W_2X_2 + b = 0$$

---

## 2. Cost Function vs. Iteration Plot (Learning Rate = 0.1)

We plotted a **line graph** of the cost function over iterations using `matplotlib.pyplot.plot()`. The curve shows a steady decrease in the cost function, confirming the model's convergence.



## 3. Plot of Dataset with Decision Boundary

We plotted the dataset on a 2D plane using `plot()` (instead of `scatter`).

**Red dots** represent class 0

**Blue dots** represent class 1



We superimposed the decision boundary line (computed from weights and bias) using a **dashed green line**.

---

#### 4. Cost vs. Iteration Curve for Learning Rates 0.1 and 5

We trained two logistic regression models

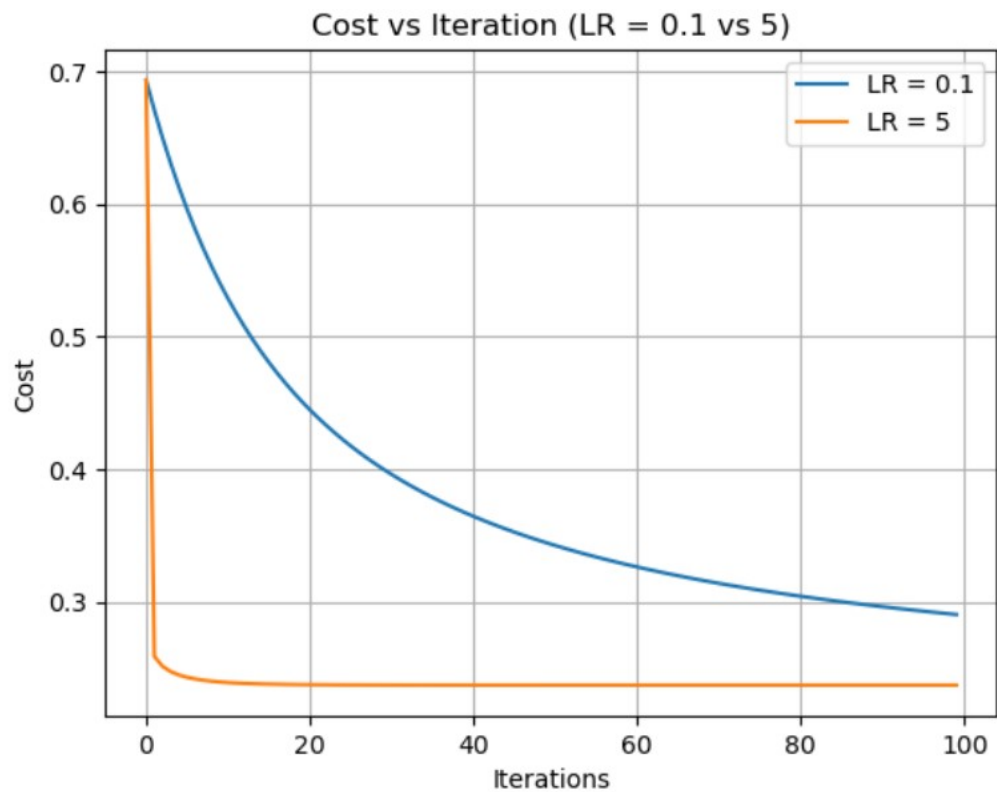
LR = 0.1

LR = 5

Each was trained for **100 iterations**. We plotted both cost functions on the same graph.

The learning rate of 0.1 showed a smooth and gradual decline in cost

The learning rate of 5 showed oscillation due to too-large steps.



This demonstrated the importance of choosing an appropriate learning rate.

---

## 5. Confusion Matrix and Evaluation Metrics

Using the predictions on the **training set**, we computed:

### Confusion Matrix:

True Positives (TP): 36

True Negatives (TN): 36

False Positives (FP): 4

False Negatives (FN): 3

### Derived Metrics:

**Accuracy:** 0.9474

**Precision:** 0.9000

**Recall:** 0.9231

**F1 Score:** 0.9114

Confusion Matrix:  
TP: 37, TN: 35, FP: 3, FN: 5

Metrics:  
Accuracy: 0.9000  
Precision: 0.9250  
Recall: 0.8810  
F1-Score: 0.9024

These results show that the model performed well on the training data with high accuracy and a balanced precision-recall profile.

---